

Local AI Setup: My Process (Frontend / E4)

What I set up

I got a local AI chat running on my laptop that can read documents and answer questions about them, all offline, nothing leaves the machine. It uses two things:

- **Ollama** - runs the AI model on my computer
- **Open WebUI** - the chat interface you actually type into

Step 1: Ollama

- Downloaded Ollama from ollama.com and installed it.
- Downloaded a model to run with this command:
`ollama pull llama3.1:8b`

Step 2: Open WebUI

Installed it with:

```
pip install open-webuiopen-webui serve
```

-
- Opened it in my browser at <http://localhost:3000> and made a local account.
- It automatically picked up the Ollama model, so I could select it in the chat.

Step 3: Testing it with a document

- Uploaded a sample PDF (an IAEA nuclear security doc, about 76 pages).
- Asked it: "Summarize this document in 5 bullet points."
- It gave an accurate summary and showed "Retrieved 1 source" with a clickable source, so I know it's actually reading the document and citing it, not just making stuff up.

Where this gets us

- We have a working chat interface running locally
- It answers questions from uploaded documents (RAG)
- It shows clickable citations